

Inferential Statistics and Predictive Analysis

Arpan Daniel Frank [RA2211047010035]

21AIC401T: Customer Churn Prediction - Case Study

Task 1: Data Preparation and Introduction

1.1 Load Libraries and Dataset

This first step involves importing the necessary Python libraries for data analysis and visualization (`pandas` , `matplotlib` , `seaborn`).

[cite_start]We will then load the **Telco Customer Churn dataset** downloaded from Kaggle[cite: 24]. [cite_start]We will perform three initial checks as required by the assignment[cite: 31, 32]:

1. `.head()` : View the first few rows to understand the columns and data structure.
2. `.info()` : Check for any missing (null) values and see the data types (e.g., object, int, float) for each column.
3. `.duplicated().sum()` : Check if there are any duplicate rows in the dataset.

```
In [1]: # --- 1. Load Libraries ---
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Set a visual style for plots
sns.set(style="whitegrid")

# --- 2. Load Data ---
# Make sure you have the 'WA_Fn-UseC_-Telco-Customer-Churn.csv' file
# in the same directory as your notebook, or provide the full path.
try:
    df = pd.read_csv('WA_Fn-UseC_-Telco-Customer-Churn.csv')

    # --- 3. Initial Inspection (Data Description & Cleaning) [cite: 31, 32] ---
    print("--- 1. First 5 Rows of the Dataset ---")
    display(df.head())

    print("\n--- 2. Dataset Info (Checking for Missing Values & Data Types) ---")
    df.info()

    print("\n--- 3. Checking for Duplicate Rows ---")
    duplicate_count = df.duplicated().sum()
    print(f"Total duplicate rows found: {duplicate_count}")

except FileNotFoundError:
    print("Error: The file 'WA_Fn-UseC_-Telco-Customer-Churn.csv' was not found.")
    print("Please download it from Kaggle and place it in the correct directory.")

--- 1. First 5 Rows of the Dataset ---
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	7590-VHVEG	Female	0	Yes	No	1	No	N													
1	5575-GNVDE	Male	0	No	No	34	Yes														
2	3668-QPYBK	Male	0	No	No	2	Yes														
3	7795-CFOCW	Male	0	No	No	45	No	N													
4	9237-HQITU	Female	0	No	No	2	Yes														

5 rows × 21 columns

```

--- 2. Dataset Info (Checking for Missing Values & Data Types) ---
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   customerID        7043 non-null   object 
 1   gender             7043 non-null   object 
 2   SeniorCitizen      7043 non-null   int64  
 3   Partner            7043 non-null   object 
 4   Dependents         7043 non-null   object 
 5   tenure             7043 non-null   int64  
 6   PhoneService       7043 non-null   object 
 7   MultipleLines      7043 non-null   object 
 8   InternetService    7043 non-null   object 
 9   OnlineSecurity     7043 non-null   object 
 10  OnlineBackup       7043 non-null   object 
 11  DeviceProtection   7043 non-null   object 
 12  TechSupport        7043 non-null   object 
 13  StreamingTV        7043 non-null   object 
 14  StreamingMovies    7043 non-null   object 
 15  Contract           7043 non-null   object 
 16  PaperlessBilling   7043 non-null   object 
 17  PaymentMethod      7043 non-null   object 
 18  MonthlyCharges    7043 non-null   float64
 19  TotalCharges       7043 non-null   object 
 20  Churn              7043 non-null   object 
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

--- 3. Checking for Duplicate Rows ---

Total duplicate rows found: 0

1.2 Data Cleaning & Target Variable Analysis

Based on the `.info()` output from the previous step, we observed two things:

1. The `TotalCharges` column is an `object` type, not a number. It needs to be converted to a numeric format. Any values that cannot be converted (like empty strings) will become `NaN` (missing values).
2. The `customerID` column is a unique identifier for each customer and is not a useful feature for prediction. We will drop it.

[cite_start]After conversion, we will handle the new missing values by simply dropping the rows, as it's the most straightforward cleaning method.

[cite_start]Finally, we will define our target variable `Churn` [cite_start]and create our first visualization to see the class distribution (i.e., how many customers churned vs. stayed). This is critical for understanding if our dataset is "imbalanced."

```
In [2]: # --- 1. Data Cleaning: Convert 'TotalCharges' to numeric ---
# 'errors=coerce' will turn any non-numeric values (like empty strings) into NaN
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')

# --- 2. Data Cleaning: Handle Missing Values ---
# Check how many missing values were created
print(f"Missing values in 'TotalCharges' after conversion: {df['TotalCharges'].isnull().sum()}

# Drop rows with any missing values
df.dropna(inplace=True)

# --- 3. Feature Engineering: Drop 'customerID' ---
# Drop the customerID column as it's not a predictive feature
df = df.drop('customerID', axis=1)

# --- 4. Define Target and Predictor Variables (EDA) ---
# Define our target variable
target_variable = 'Churn'

# --- 5. Conduct EDA: Visualize the Target Variable ---
print(f"\n--- Distribution of Target Variable: {target_variable} ---")

# Create a count plot to visualize the 'Churn' column
plt.figure(figsize=(7, 5))
sns.countplot(x=target_variable, data=df)
plt.title('Churn Distribution (No vs. Yes)')


# Calculate and print the percentage
churn_percentage = (df[target_variable].value_counts(normalize=True) * 100).reset_index()
churn_counts = df[target_variable].value_counts().rename('Count')
print(pd.concat([churn_counts, churn_percentage], axis=1))

# Show the plot
plt.show()
```

Missing values in 'TotalCharges' after conversion: 11

--- Distribution of Target Variable: Churn ---
Count Percentage

	Count	Percentage
No	5163	73.421502
Yes	1869	26.578498



1.3 EDA: Numerical Variable Analysis

Now we perform Exploratory Data Analysis (EDA) on the numerical features to understand their relationship with `Churn`.

We will iterate through our main numerical columns (`tenure`, `MonthlyCharges`, `TotalCharges`) and create a Kernel Density Estimate (KDE) plot for each.

These plots will help us visualize the distribution of each feature, separated by customers who churned (`Churn` = 'Yes') and those who stayed (`Churn` = 'No'). This can reveal important patterns (e.g., "Do customers with lower tenure churn more?").

```
In [3]: # --- 1. Identify Numerical Predictor Variables ---
# From our .info() earlier, we know these are our numeric features
numerical_features = ['tenure', 'MonthlyCharges', 'TotalCharges']

print("--- EDA: Analyzing Numerical Features vs. Churn ---")

# --- 2. Loop and Plot (KDE Plots) ---
for col in numerical_features:
    plt.figure(figsize=(10, 6))

    # Plot KDE for customers who did NOT churn
    sns.kdeplot(df[df['Churn'] == 'No'][col], label='Churn: No', shade=True, col

    # Plot KDE for customers who DID churn
    sns.kdeplot(df[df['Churn'] == 'Yes'][col], label='Churn: Yes', shade=True, c

    # Add titles and Labels
    plt.title(f'Distribution of {col} by Churn')
```

```

plt.xlabel(col)
plt.ylabel('Density')
plt.legend()

# Show the plot
plt.show()

```

--- EDA: Analyzing Numerical Features vs. Churn ---

C:\Users\hp pavillion\AppData\Local\Temp\ipykernel_14092\1105170387.py:12: Future Warning:


`shade` is now deprecated in favor of `fill`; setting `fill=True`. This will become an error in seaborn v0.14.0; please update your code.

 sns.kdeplot(df[df['Churn'] == 'No'][col], label='Churn: No', shade=True, color='blue')

C:\Users\hp pavillion\AppData\Local\Temp\ipykernel_14092\1105170387.py:15: Future Warning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`. This will become an error in seaborn v0.14.0; please update your code.

 sns.kdeplot(df[df['Churn'] == 'Yes'][col], label='Churn: Yes', shade=True, color='red')




```
C:\Users\hp pavillion\AppData\Local\Temp\ipykernel_14092\1105170387.py:12: Future
Warning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

sns.kdeplot(df[df['Churn'] == 'No'][col], label='Churn: No', shade=True, color
='blue')
C:\Users\hp pavillion\AppData\Local\Temp\ipykernel_14092\1105170387.py:15: Future
Warning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

sns.kdeplot(df[df['Churn'] == 'Yes'][col], label='Churn: Yes', shade=True, colo
r='red')
```




```
C:\Users\hp pavillion\AppData\Local\Temp\ipykernel_14092\1105170387.py:12: Future
Warning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

sns.kdeplot(df[df['Churn'] == 'No'][col], label='Churn: No', shade=True, color
='blue')
C:\Users\hp pavillion\AppData\Local\Temp\ipykernel_14092\1105170387.py:15: Future
Warning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

sns.kdeplot(df[df['Churn'] == 'Yes'][col], label='Churn: Yes', shade=True, colo
r='red')
```



1.4 EDA: Categorical Variable Analysis

Next, we analyze the categorical features to see their impact on churn.

We will iterate through our main categorical columns and create a `countplot` for each.

We will use the `hue='Churn'` parameter in our plots.

This will create stacked bar charts, allowing us to easily see not just the total number of customers in each category (e.g., "Month-to-month" contract), but also the proportion of those customers who churned vs. stayed. This is one of the most effective ways to find key drivers of churn.

```
In [4]: # --- 1. Identify Categorical Predictor Variables ---
# We list them manually for clarity, excluding our target 'Churn'
categorical_features = [
    'gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines',
    'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
    'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
    'PaperlessBilling', 'PaymentMethod'
]

print("--- EDA: Analyzing Categorical Features vs. Churn ---")

# --- 2. Loop and Plot (Count Plots) ---
for col in categorical_features:
    plt.figure(figsize=(10, 6))


    # Create a count plot, splitting by the 'Churn' variable
    sns.countplot(data=df, x=col, hue='Churn')



    # Add titles and labels
    plt.title(f'Churn Rate by {col}')
    plt.xlabel(col)
    plt.ylabel('Number of Customers')
```



```
# Rotate x-axis Labels if they are Long
if df[col].nunique() > 4:
    plt.xticks(rotation=30)



# Show the plot
plt.show()
```



--- EDA: Analyzing Categorical Features vs. Churn ---






















Task 2 (Part 1): Data Preprocessing for CHAID

We are now starting Task 2: Model Development using CHAID.

The CHAID algorithm requires all of its predictor variables to be categorical. Our dataset currently has a mix of numerical and categorical features.

Our plan is to:

- 1. Discretize (Bin) Numerical Features:** Convert `tenure`, `MonthlyCharges`, and `TotalCharges` into categorical bins (e.g., 4 groups based on quantiles).
- 2. Encode All Features:** Apply `LabelEncoder` to *all* predictor columns. This will convert text categories (like 'Yes', 'No', 'DSL') into integers (like 1, 0, 2), which the CHAID library requires.
- 3. Split Data:** Separate our data into a training set (to build the model) and a testing set (to evaluate it later).

```
In [5]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

# --- 1. Create a copy of the dataframe for CHAID preprocessing ---
# This ensures our original 'df' remains unchanged for other models
df_chaid = df.copy()

# --- 2. Discretize (Bin) Numerical Features ---
# We will use 'qcut' (quantile-cut) to split them into 4 equal-sized groups
numerical_features = ['tenure', 'MonthlyCharges', 'TotalCharges']
for col in numerical_features:
    df_chaid[col] = pd.qcut(df_chaid[col], q=4, duplicates='drop').astype(str)

# --- 3. Encode All Features (Predictors and Target) ---
# Separate predictors (X) and target (y)
```

```

X = df_chaid.drop('Churn', axis=1)
y = df_chaid['Churn']

# Apply LabelEncoder to all predictor columns in X
X_encoded = X.apply(LabelEncoder().fit_transform)

# Apply LabelEncoder to the target variable y
y_encoded = LabelEncoder().fit_transform(y)
# (In 'y', 'No' will become 0 and 'Yes' will become 1)

print("--- Data after Encoding (First 5 Rows) ---")
display(X_encoded.head())

# --- 4. Split Data into Training and Testing Sets ---
# We'll use 80% for training and 20% for testing
X_train, X_test, y_train, y_test = train_test_split(
    X_encoded,
    y_encoded,
    test_size=0.2,
    random_state=42
)

print(f"\nData split successfully:")
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")

```

--- Data after Encoding (First 5 Rows) ---

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	Inte
0	0	0	1	0	0	0	0	1
1	1	0	0	0	1	1	1	0
2	1	0	0	0	0	1	1	0
3	1	0	0	0	1	0	0	1
4	0	0	0	0	0	1	1	0

◀ ⟳

Data split successfully:
X_train shape: (5625, 19)
X_test shape: (1407, 19)
y_train shape: (5625,)
y_test shape: (1407,)

Task 2 (Part 2): Build CHAID Model & Extract Rules

Now that our data is preprocessed, we can build the CHAID model.

1. **Install Library:** We will first install the `CHAID` library, as it is not a default part of `scikit-learn`.
2. **Build Tree:** We will instantiate the `Tree` object from the library and fit it on our training data (`X_train`, `y_train`). We'll set a `max_depth` (e.g., 5) to keep the tree from getting too large and to make the rules easier to interpret.

3. **Extract Rules:** The key output of this model is the set of human-readable decision rules. We will print the resulting tree structure, which shows the "IF-THEN" logic that the model discovered. [cite_start]This directly addresses the assignment's requirement to "extract decision rules"[cite: 36].

```
In [8]: # --- 1. Define Model Inputs ---
# We use 'df_chaid', the dataframe with binned numerical features
# and original string categories (e.g., 'Yes', 'No')
from CHAID import Tree

predictor_columns = [col for col in df_chaid.columns if col != 'Churn']
target_column = 'Churn'
column_types = {col: 'nominal' for col in predictor_columns}

# --- 2. Build the CHAID Tree (on the full dataset) ---
# The goal here is rule discovery, not prediction, so
# we use the full dataset to find all possible patterns.
print("--- Building CHAID Tree (max_depth=5) ---")
chaid_tree = Tree.from_pandas_df(
    df_chaid,                      # The dataframe with raw categories
    column_types,                  # The dictionary mapping predictors to 'nominal'
    target_column,                 # The name of the target column
    max_depth=5                    # Set the maximum depth to 5
)

# --- 3. Print the Decision Rules ---
print("\n--- CHAID Decision Rules (for Task 2 Report) ---")
chaid_tree.print_tree()

print("\nTask 2 complete. Please copy the rules above for your report.")
```

```
c:\Users\hp pavillion\AppData\Local\Programs\Python\Python310\lib\site-packages\CHAID\graph.py:13: UserWarning: Imports of optional packages needed to generate graphs failed. Please install with the "graph" option.
  warnings.warn(UserWarning('Imports of optional packages needed to generate graphs failed. Please install with the "graph" option.'))
```

```

--- Building CHAID Tree (max_depth=5) ---

--- CHAID Decision Rules (for Task 2 Report) ---
([], {'No': 5163.0, 'Yes': 1869.0}, (Contract, p=7.326182186265472e-257, score=11
79.5458287339445, groups=[['Month-to-month'], ['One year'], ['Two year']]), dof=
2))
|-- ([('Month-to-month'), {'No': 2220.0, 'Yes': 1655.0}, (InternetService, p=8.853
102500484912e-66, score=299.5796963499884, groups=[['DSL'], ['Fiber optic'], ['N
o']]), dof=2))
|  |-- ([('DSL'), {'No': 829.0, 'Yes': 394.0}, (TotalCharges, p=6.320879510155006
e-25, score=111.44153792675088, groups=[['(1397.475, 3794.738]', '(3794.738, 868
4.8]'], ['(18.799, 401.45]'], ['(401.45, 1397.475]']), dof=2))
|  |  |-- ([('1397.475, 3794.738]', '(3794.738, 8684.8]'], {'No': 229.0, 'Yes':
38.0}, (PhoneService, p=0.001875614417503133, score=9.667450235354163, groups=
[['No'], ['Yes']]), dof=1))
|  |  |  |-- ([('No'), {'No': 58.0, 'Yes': 19.0}, <Invalid Chaid Split> - the n
ode only contains single category respondents)
|  |  |  |  +-- ([('Yes'), {'No': 171.0, 'Yes': 19.0}, <Invalid Chaid Split> - p-v
alue greater than alpha merge]
|  |  |  |  +-- ([('18.799, 401.45]'], {'No': 294.0, 'Yes': 263.0}, (PaymentMethod, p
=0.0014564321619320209, score=10.132936247419636, groups=[['Bank transfer (automa
tic)', 'Credit card (automatic)', 'Mailed check'], ['Electronic check']]), dof=
1))
|  |  |  |  +-- ([('Bank transfer (automatic)', 'Credit card (automatic)', 'Mailed
check'), {'No': 189.0, 'Yes': 134.0}, (OnlineSecurity, p=0.007403242048135513, sc
ore=7.172430147442203, groups=[['No'], ['Yes']]), dof=1])
|  |  |  |  |-- ([('No'), {'No': 137.0, 'Yes': 114.0}, <Invalid Chaid Split> - 
the max depth has been reached)
|  |  |  |  |  +-- ([('Yes'), {'No': 52.0, 'Yes': 20.0}, <Invalid Chaid Split> - 
the max depth has been reached)
|  |  |  |  |  +-- ([('Electronic check'), {'No': 105.0, 'Yes': 129.0}, (SeniorCitize
n, p=0.011972888128015327, score=6.31487951011844, groups=[[0], [1]]), dof=1])
|  |  |  |  |  |-- ([0], {'No': 95.0, 'Yes': 101.0}, <Invalid Chaid Split> - the
max depth has been reached)
|  |  |  |  |  |  +-- ([1], {'No': 10.0, 'Yes': 28.0}, <Invalid Chaid Split> - the
max depth has been reached)
|  |  |  |  |  |  +-- ([('401.45, 1397.475']], {'No': 306.0, 'Yes': 93.0}, (OnlineBackup, p
=0.0005570716992938892, score=11.914216275767004, groups=[['No'], ['Yes']]), dof=
1))
|  |  |  |  |  |  |-- ([('No'), {'No': 203.0, 'Yes': 79.0}, (SeniorCitizen, p=0.00455809
02310967, score=8.046927195146937, groups=[[0], [1]]), dof=1])
|  |  |  |  |  |  |  +-- ([0], {'No': 188.0, 'Yes': 64.0}, <Invalid Chaid Split> - the
max depth has been reached)
|  |  |  |  |  |  |  +-- ([1], {'No': 15.0, 'Yes': 15.0}, <Invalid Chaid Split> - the
minimum parent node size threshold has been reached)
|  |  |  |  |  |  |  +-- ([('Yes'), {'No': 103.0, 'Yes': 14.0}, <Invalid Chaid Split> - the
node only contains single category respondents)
|  |  |  |  |  |  |  |-- ([('Fiber optic'), {'No': 966.0, 'Yes': 1162.0}, (tenure, p=6.085068284360
564e-41, score=190.00621994037962, groups=[['(0.999, 9.0]'], ['(29.0, 55.0]'],
['(55.0, 72.0]'], ['(9.0, 29.0]']), dof=3))
|  |  |  |  |  |  |  |  +-- ([('0.999, 9.0)'), {'No': 226.0, 'Yes': 559.0}, (OnlineSecurity, p=0.
00012134122323093427, score=14.7717271650391, groups=[['No'], ['Yes']]), dof=1)
|  |  |  |  |  |  |  |  |-- ([('No'), {'No': 195.0, 'Yes': 528.0}, (TotalCharges, p=0.00091645
62189801672, score=10.989174497523898, groups=[['(18.799, 401.45]'], ['(401.45, 1
397.475]']), dof=1))
|  |  |  |  |  |  |  |  |  |-- ([('18.799, 401.45]'], {'No': 116.0, 'Yes': 382.0}, <Invalid
Chaid Split> - the max depth has been reached)
|  |  |  |  |  |  |  |  |  |  +-- ([('401.45, 1397.475']], {'No': 79.0, 'Yes': 146.0}, <Invalid
Chaid Split> - the max depth has been reached)
|  |  |  |  |  |  |  |  |  |  +-- ([('Yes'), {'No': 31.0, 'Yes': 31.0}, <Invalid Chaid Split> - p-va

```

```

lue greater than alpha merge)
|   |   |-- ([(29.0, 55.0)], {'No': 299.0, 'Yes': 204.0}, (PaymentMethod, p=0.0
008613048375408288, score=11.104242024179086, groups=[['Bank transfer (automati
c)', 'Credit card (automatic)', 'Mailed check'], ['Electronic check']]), dof=1))
|   |   |   |-- (['Bank transfer (automatic)', 'Credit card (automatic)', 'Mailed
check'], {'No': 158.0, 'Yes': 77.0}, (SeniorCitizen, p=0.015368010211948608, scor
e=5.873780991394172, groups=[[0], [1]]), dof=1))
|   |   |   |   |-- ([0], {'No': 124.0, 'Yes': 49.0}, <Invalid Chaid Split> - the
max depth has been reached)
|   |   |   |   +- ([1], {'No': 34.0, 'Yes': 28.0}, <Invalid Chaid Split> - the
max depth has been reached)
|   |   |   +- (['Electronic check'], {'No': 141.0, 'Yes': 127.0}, (OnlineBacku
p, p=0.023531903042689425, score=5.128806526025489, groups=[['No'], ['Yes']]), do
f=1))
|   |   |       |-- (['No'], {'No': 66.0, 'Yes': 77.0}, <Invalid Chaid Split> - t
he max depth has been reached)
|   |   |       +- (['Yes'], {'No': 75.0, 'Yes': 50.0}, <Invalid Chaid Split> -
the max depth has been reached)
|   |   |   |-- ([(55.0, 72.0)], {'No': 121.0, 'Yes': 37.0}, (StreamingTV, p=0.0031
252887067114853, score=8.732842651714124, groups=[['No'], ['Yes']]), dof=1))
|   |   |   |   |-- (['No'], {'No': 52.0, 'Yes': 6.0}, <Invalid Chaid Split> - the no
de only contains single category respondents)
|   |   |   |   +- (['Yes'], {'No': 69.0, 'Yes': 31.0}, (TechSupport, p=0.0087230699
42273122, score=6.878664906790146, groups=[['No'], ['Yes']]), dof=1))
|   |   |       |-- (['No'], {'No': 42.0, 'Yes': 27.0}, <Invalid Chaid Split> - t
he max depth has been reached)
|   |   |       +- (['Yes'], {'No': 27.0, 'Yes': 4.0}, <Invalid Chaid Split> - t
he max depth has been reached)
|   |   +- ([(9.0, 29.0)], {'No': 320.0, 'Yes': 362.0}, (PaymentMethod, p=1.85
75843171978175e-05, score=18.330003854262397, groups=[['Bank transfer (automati
c)', 'Mailed check', 'Credit card (automatic)'], ['Electronic check']]), dof=1))
|   |       |-- (['Bank transfer (automatic)', 'Mailed check', 'Credit card (auto
matic)'], {'No': 153.0, 'Yes': 115.0}, (StreamingMovies, p=7.150171298599028e-06,
score=20.15263410455226, groups=[['No'], ['Yes']]), dof=1))
|   |       |   |-- (['No'], {'No': 107.0, 'Yes': 49.0}, <Invalid Chaid Split> -
the max depth has been reached)
|   |       |   +- (['Yes'], {'No': 46.0, 'Yes': 66.0}, <Invalid Chaid Split> -
the max depth has been reached)
|   |       +- (['Electronic check'], {'No': 167.0, 'Yes': 247.0}, (MultipleLine
s, p=0.0014082648476866737, score=10.194940754992746, groups=[['No'], ['Yes']]),
dof=1))
|   |           |-- (['No'], {'No': 83.0, 'Yes': 84.0}, <Invalid Chaid Split> - t
he max depth has been reached)
|   |           +- (['Yes'], {'No': 84.0, 'Yes': 163.0}, <Invalid Chaid Split> -
the max depth has been reached)
|   |   +- (['No'], {'No': 425.0, 'Yes': 99.0}, (tenure, p=7.939471301056829e-06,
score=19.952388453619644, groups=[['(0.999, 9.0]', '(55.0, 72.0]', '(29.0, 55.
0]', '(9.0, 29.0]']), dof=1))
|   |       |-- ([(0.999, 9.0]', '(55.0, 72.0]'), {'No': 260.0, 'Yes': 84.0}, <Inval
id Chaid Split> - p-value greater than alpha merge)
|   |       +- ([(29.0, 55.0]', '(9.0, 29.0)'), {'No': 165.0, 'Yes': 15.0}, (TotalC
harges, p=0.037915871054459316, score=4.308792253043124, groups=[['(18.799, 401.4
5]', '(401.45, 1397.475]']), dof=1))
|   |           |-- ([(18.799, 401.45]', '(401.45, 1397.475)'), {'No': 86.0, 'Yes': 12.0}, <Invalid Chaid
Split> - the node only contains single category respondents)
|   |           +- ([(401.45, 1397.475)], {'No': 79.0, 'Yes': 3.0}, <Invalid Chaid
Split> - the node only contains single category respondents)
|   |   |-- (['One year'], {'No': 1306.0, 'Yes': 166.0}, (StreamingMovies, p=2.2256538457
43632e-18, score=81.29296187637681, groups=[['No'], ['No internet service'], ['Ye
s']]), dof=2))

```

```

|   |-- (['No'], {'No': 418.0, 'Yes': 29.0}, (PaymentMethod, p=0.0002788586453788
878, score=13.207314577437081, groups=[['Bank transfer (automatic)', 'Credit card
(automatic)', 'Mailed check'], ['Electronic check']]), dof=1)
|   |   |-- (['Bank transfer (automatic)', 'Credit card (automatic)', 'Mailed che
ck'], {'No': 336.0, 'Yes': 15.0}, (TotalCharges, p=0.009704976340340725, score=6.
6882621951219505, groups=[['(1397.475, 3794.738]', '(18.799, 401.45]', '(401.45,
1397.475]'], [(3794.738, 8684.8]]], dof=1)
|   |   |   |-- ([(1397.475, 3794.738]', '(18.799, 401.45]', '(401.45, 1397.47
5]'], {'No': 231.0, 'Yes': 15.0}, <Invalid Chaid Split> - p-value greater than al
pha merge)
|   |   |   +- ([(3794.738, 8684.8]], {'No': 105.0, 'Yes': 0}, <Invalid Chaid
Split> - the node only contains single category respondents)
|   |   +- (['Electronic check'], {'No': 82.0, 'Yes': 14.0}, <Invalid Chaid Spli
t> - p-value greater than alpha merge)
|   |-- (['No internet service'], {'No': 354.0, 'Yes': 9.0}, (PaymentMethod, p=0.
03731426562874419, score=4.336010257794574, groups=[['Bank transfer (automatic)'],
'Electronic check'], ['Credit card (automatic)', 'Mailed check']]), dof=1)
|   |   |-- (['Bank transfer (automatic)', 'Electronic check'], {'No': 118.0, 'Ye
s': 6.0}, <Invalid Chaid Split> - p-value greater than alpha merge)
|   |   +- (['Credit card (automatic)', 'Mailed check'], {'No': 236.0, 'Yes': 3.
0}, (tenure, p=0.0006516610816516224, score=14.67197188921042, groups=[['(0.999,
9.0]', '(29.0, 55.0]', '(55.0, 72.0]', '(9.0, 29.0]']]), dof=2)
|   |   |   |-- ([(0.999, 9.0]', '(29.0, 55.0]', '(55.0, 72.0]', '(9.0, 29.0]'],
dof=2)
|   |   |   +- ([(0.999, 9.0]', {'No': 38.0, 'Yes': 3.0}, <Invalid Chaid Split
> - the node only contains single category respondents)
|   |   |   |-- ([(29.0, 55.0]', '(55.0, 72.0]', {'No': 88.0, 'Yes': 0}, <Inval
id Chaid Split> - the node only contains single category respondents)
|   |   |   +- ([(9.0, 29.0]', {'No': 110.0, 'Yes': 0}, <Invalid Chaid Split
> - the node only contains single category respondents)
|   |   +- (['Yes'], {'No': 534.0, 'Yes': 128.0}, (InternetService, p=0.001185609188
2215625, score=10.51263738147573, groups=[['DSL'], ['Fiber optic']]), dof=1)
|   |   |-- (['DSL'], {'No': 225.0, 'Yes': 34.0}, <Invalid Chaid Split> - p-value
greater than alpha merge)
|   |   +- (['Fiber optic'], {'No': 309.0, 'Yes': 94.0}, (StreamingTV, p=0.03093
7030512007014, score=4.656453438543821, groups=[['No'], ['Yes']]), dof=1)
|   |   |-- (['No'], {'No': 59.0, 'Yes': 9.0}, (DeviceProtection, p=0.0158956
60876330597, score=5.814379468616755, groups=[['No'], ['Yes']]), dof=1)
|   |   |   |-- (['No'], {'No': 27.0, 'Yes': 8.0}, <Invalid Chaid Split> - th
e max depth has been reached)
|   |   |   +- (['Yes'], {'No': 32.0, 'Yes': 1.0}, <Invalid Chaid Split> - t
he max depth has been reached)
|   |   +- (['Yes'], {'No': 250.0, 'Yes': 85.0}, <Invalid Chaid Split> - p-v
alue greater than alpha merge)
+- (['Two year'], {'No': 1637.0, 'Yes': 48.0}, (InternetService, p=2.74362621097
35596e-10, score=39.84908551133631, groups=[['DSL', 'No'], ['Fiber optic']]), dof
=1))
|-- (['DSL', 'No'], {'No': 1239.0, 'Yes': 17.0}, (SeniorCitizen, p=0.00187404
91036021263, score=9.668984114864745, groups=[[0], [1]]), dof=1)
|   |-- ([0], {'No': 1169.0, 'Yes': 13.0}, (PaymentMethod, p=0.00021216545788
49526, score=13.720154144172987, groups=[['Bank transfer (automatic)', 'Credit ca
rd (automatic)', 'Mailed check'], ['Electronic check']]), dof=1)
|   |   |-- (['Bank transfer (automatic)', 'Credit card (automatic)', 'Mailed che
ck'], {'No': 1100.0, 'Yes': 9.0}, <Invalid Chaid Split> - p-value greater than al
pha merge)
|   |   +- (['Electronic check'], {'No': 69.0, 'Yes': 4.0}, <Invalid Chaid S
plit> - p-value greater than alpha merge)
|   |   +- ([1], {'No': 70.0, 'Yes': 4.0}, (StreamingTV, p=0.03462664031644337,
score=4.463492063492064, groups=[['No', 'No internet service'], ['Yes']]), dof=
1))
|   |   |-- (['No', 'No internet service'], {'No': 38.0, 'Yes': 0}, <Invalid
Chaid Split> - the node only contains single category respondents)

```

```

|      +-+ ([ 'Yes' ], { 'No': 32.0, 'Yes': 4.0 }, <Invalid Chaid Split> - p-val
ue greater than alpha merge)
+-+ ([ 'Fiber optic' ], { 'No': 398.0, 'Yes': 31.0 }, (PaymentMethod, p=0.0321835
3990434876, score=4.58869337798977, groups=[ [ 'Bank transfer (automatic)', 'Electr
onic check' ], [ 'Credit card (automatic)', 'Mailed check' ] ] ), dof=1)
|-- ([ 'Bank transfer (automatic)', 'Electronic check' ], { 'No': 230.0, 'Ye
s': 24.0 }, <Invalid Chaid Split> - p-value greater than alpha merge)
+-+ ([ 'Credit card (automatic)', 'Mailed check' ], { 'No': 168.0, 'Yes': 7.
0 }, (Partner, p=0.031601199059749494, score=4.6200021235287, groups=[ [ 'No' ], [ 'Ye
s' ] ] ), dof=1)
|-- ([ 'No' ], { 'No': 37.0, 'Yes': 4.0 }, <Invalid Chaid Split> - p-valu
e greater than alpha merge)
+-+ ([ 'Yes' ], { 'No': 131.0, 'Yes': 3.0 }, <Invalid Chaid Split> - p-va
lue greater than alpha merge)

```

Task 2 complete. Please copy the rules above for your report.

Task 3 (Part 2): Build & Train Models (Decision Tree and Logistic Regression)

Now that our data is fully numeric and split, we can build the two models for our comparison.

1. **Model 1: Decision Tree Classifier:** This is our functional replacement for CHAID, as it's a tree-based model that can be used for prediction.
2. [cite_start]**Model 2: Logistic Regression:** This is the second model required by the assignment[cite: 44].

We will:

1. Import, initialize, and train both models on our `X_train` data.
2. Get two types of predictions from each model on the `X_test` data:
 - **Class Predictions (`.predict()`):** The final 'Yes' (1) or 'No' (0) guess. Used for calculating **Accuracy**.
 - **Probability Predictions (`.predict_proba()`):** The model's confidence (e.g., "70% chance of churn"). This is required to calculate **ROC-AUC** and draw **Lift/Gains charts**.

```
In [9]: from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler

# --- 1. Scale the Data (Important for Logistic Regression) ---

# Logistic Regression performs better when features are on a similar scale.
# Decision Trees don't require scaling, but it doesn't hurt them.
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# --- 2. Build and Train Model 1: Decision Tree ---
print("--- Training Decision Tree ---")
# We set max_depth=5 to match our CHAID tree and prevent overfitting
```

```

dt_model = DecisionTreeClassifier(max_depth=5, random_state=42)
dt_model.fit(X_train_scaled, y_train)

# --- 3. Build and Train Model 2: Logistic Regression ---
print("--- Training Logistic Regression ---")
# 'max_iter' is increased to 1000 to ensure the model converges
lr_model = LogisticRegression(max_iter=1000, random_state=42)
lr_model.fit(X_train_scaled, y_train)

# --- 4. Get Predictions from Both Models ---
print("--- Generating predictions ---")

# Decision Tree predictions
y_pred_dt = dt_model.predict(X_test_scaled)
y_proba_dt = dt_model.predict_proba(X_test_scaled)[:, 1] # Get probability of ch

# Logistic Regression predictions
y_pred_lr = lr_model.predict(X_test_scaled)
y_proba_lr = lr_model.predict_proba(X_test_scaled)[:, 1] # Get probability of ch

print("Models trained and predictions generated successfully.")

```

--- Training Decision Tree ---
--- Training Logistic Regression ---
--- Generating predictions ---
Models trained and predictions generated successfully.

Task 3 (Part 3): Model Comparison & Evaluation (Metrics)

This is the core of Task 3. We will now use the predictions we generated in the previous step to evaluate and compare our two models.

[cite_start]As required by the assignment, we will use the following metrics:

1. **Accuracy**: The percentage of correct predictions.
2. **ROC-AUC Score**: A measure of how well the model can distinguish between churners and non-churners.
3. **ROC Curve**: A visual plot of the model's performance.
4. **Gains and Lift Charts**: These are used to show the "lift" or "gain" from using the model, which is a key business metric.

We will need to install the `scikit-plot` library to easily create the Gains and Lift charts.

```

In [11]: # --- 1. Import all necessary metrics libraries ---
from sklearn.metrics import accuracy_score, roc_auc_score, roc_curve
import scikitplot as skplt
import matplotlib.pyplot as plt

# --- 2. Calculate Accuracy ---
acc_dt = accuracy_score(y_test, y_pred_dt)
acc_lr = accuracy_score(y_test, y_pred_lr)

print("--- Model Comparison ---")
print(f"Decision Tree Accuracy: {acc_dt:.4f}")

```

```

print(f"Logistic Regression Accuracy: {acc_lr:.4f}")

# --- 3. Calculate ROC-AUC Score ---
auc_dt = roc_auc_score(y_test, y_proba_dt)
auc_lr = roc_auc_score(y_test, y_proba_lr)

print(f"\nDecision Tree ROC-AUC Score: {auc_dt:.4f}")
print(f"Logistic Regression ROC-AUC Score: {auc_lr:.4f}")

# --- 4. Plot ROC Curve ---
# (Removed the 'Random Guess' Line as requested)
fpr_dt, tpr_dt, _ = roc_curve(y_test, y_proba_dt)
fpr_lr, tpr_lr, _ = roc_curve(y_test, y_proba_lr)

plt.figure(figsize=(10, 7))
plt.plot(fpr_dt, tpr_dt, label=f'Decision Tree (AUC = {auc_dt:.4f})')
plt.plot(fpr_lr, tpr_lr, label=f'Logistic Regression (AUC = {auc_lr:.4f})')
# plt.plot([0, 1], [0, 1], 'k--', Label='Random Guess') # This line is now removed
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve Comparison')
plt.legend()
plt.show()

# --- 5. Plot Gains and Lift Charts (Corrected Method) ---
# The previous ValueError is a bug. We will fix it by
# plotting each model's chart in its own figure.

# Get the 2D probability arrays for both models
probas_dt = dt_model.predict_proba(X_test_scaled)
probas_lr = lr_model.predict_proba(X_test_scaled)

print("\n--- Gains Charts ---")
skplt.metrics.plot_cumulative_gain(y_test, probas_dt, title='Cumulative Gains Chart')
plt.show()
skplt.metrics.plot_cumulative_gain(y_test, probas_lr, title='Cumulative Gains Chart')
plt.show()

print("\n--- Lift Charts ---")
skplt.metrics.plot_lift_curve(y_test, probas_dt, title='Lift Curve (Decision Tree)')
plt.show()
skplt.metrics.plot_lift_curve(y_test, probas_lr, title='Lift Curve (Logistic Regression)')
plt.show()

print("\nTask 3 (Evaluation) complete.")

```


--- Model Comparison ---

Decision Tree Accuracy: 0.7669


Logistic Regression Accuracy: 0.7676


Decision Tree ROC-AUC Score: 0.8030

Logistic Regression ROC-AUC Score: 0.8105





--- Gains Charts ---





--- Lift Charts ---





Task 3 (Evaluation) complete.

Task 4: Model Deployment and Updating

This task requires us to *explain* the process for deploying and updating our churn model. We do not need to build a full application.

4.1 Model Deployment Process

The first step in deployment is to save (or "pickle") our best-performing model. Based on our ROC-AUC scores, the **Logistic Regression** model was our best model.

We will use the `joblib` library to save our trained `lr_model` and the `scaler` object (which is required for preprocessing new data) to disk.

My Deployment Plan (for the report):

1. **Save Model:** Save the trained `lr_model` and `scaler` objects as `.joblib` files.
2. **Build API:** Create a simple web API (using a framework like **Flask** or **FastAPI**).
3. **Load Model:** The API will load the saved `lr_model` and `scaler` files when it starts.
4. **Create Endpoint:** The API will have an endpoint (e.g., `/predict`) that accepts new customer data (as JSON).
5. **Process & Predict:** The API will:
 - Receive the new data.
 - Apply the same One-Hot Encoding and scaling (using the saved `scaler`).

- Pass the processed data to the `lr_model.predict_proba()` method.
- Return a JSON response with the churn probability.

4.2 Model Updating Process

Models grow "stale" as new customer data and behaviors emerge.

My Updating Plan (for the report):

1. **Collect New Data:** Set up a process to automatically collect and label new customer data (e.g., customers who churned in the last month).
2. **Schedule Retraining:** Create a scheduled script (e.g., a monthly **CRON job** or **Airflow DAG**) that re-runs this entire notebook (from preprocessing to training) on the *new, combined* dataset.
3. **Validate & Deploy:** After retraining, the script would automatically test the new model. If its performance (e.g., ROC-AUC) on a test set is better than the old model, the script will automatically save it, replacing the old `.joblib` file in the API.

```
In [12]: import joblib

# --- 1. Save the Logistic Regression Model ---
# This was our best model
model_filename = 'churn_model_lr.joblib'
joblib.dump(lr_model, model_filename)

# --- 2. Save the Scaler ---
# We MUST save the scaler, as new data needs to be scaled
# in the exact same way as the training data.
scaler_filename = 'data_scaler.joblib'
joblib.dump(scaler, scaler_filename)

print(f"--- Task 4: Deployment ---")
print(f"Model saved as: {model_filename}")
print(f"Scaler saved as: {scaler_filename}")
print("\nThese two files are all you need for your API.")
print("You can now write the descriptive part of Task 4 in your report.")
```

```
--- Task 4: Deployment ---
Model saved as: churn_model_lr.joblib
Scaler saved as: data_scaler.joblib
```

These two files are all you need for your API.
 You can now write the descriptive part of Task 4 in your report.